

Regression: Hands-on Session

Ditty Mathew

Machine Learning Camp

22nd June 2018

- NumPy stands for **Numerical Python**
- To import numpy in python
 - `import numpy`

array: N-dimensional array; collection of items of same type

array: N-dimensional array; collection of items of same type
`import numpy as np`

array: N-dimensional array; collection of items of same type

```
import numpy as np  
a= np.array([1,2,3])
```

array: N-dimensional array; collection of items of same type

```
import numpy as np  
a= np.array([1,2,3])  
print a
```

array: N-dimensional array; collection of items of same type

```
import numpy as np
a= np.array([1,2,3])
print a
print a.shape
```

array: N-dimensional array; collection of items of same type

```
import numpy as np
a= np.array([1,2,3])
print a
print a.shape
```

- array of more than one dimensions

```
a = np.array([[1, 2], [3, 4]])
print a
```


NumPy also provides a reshape function to resize an array.

```
a = np.array([[1,2,3],[4,5,6]])  
b = a.reshape(3,2)  
print b
```

To append a column to a numpy array

```
a = np.array([[1,2,3],[4,5,6]])  
np.column_stack((a,[7,8]))
```

To return a new array of specified size, filled with zeros

```
np.zeros((5,2))
```

To return a new array of specified size, filled with ones

```
np.ones((5,2))
```

```
a = np.array([[1,2,3],[4,5,6]])
```

To fetch value of i^{th} row and j^{th} column

```
a[i,j]
```

To fetch all values in j^{th} column

- `a[:,j]`

```
a = np.array([[1,2,3],[4,5,6]])
```

To fetch value of i^{th} row and j^{th} column

```
a[i,j]
```

To fetch all values in j^{th} column

- `a[:,j]`

To fetch all values in i^{th} row

- `a[i,:]`

To multiply two matrices

```
x=np.array([[1,2],[3,4]])  
y=np.array([[1,2,3],[3,4,5]])  
np.dot(x,y)
```

To multiply two matrices

```
x=np.array([[1,2],[3,4]])  
y=np.array([[1,2,3],[3,4,5]])  
np.dot(x,y)
```

To find a transpose of a matrix

```
y=np.array([[1,2,3],[3,4,5]])  
y.transpose()
```

To multiply two matrices

```
x=np.array([[1,2],[3,4]])  
y=np.array([[1,2,3],[3,4,5]])  
np.dot(x,y)
```

To find a transpose of a matrix

```
y=np.array([[1,2,3],[3,4,5]])  
y.transpose()
```

To find inverse of a matrix

```
y=np.array([[1,2,3],[3,4,5]])  
np.linalg.inv(y)
```


- To load data from file

```
data = np.loadtxt(open("data.csv", "rb"), delimiter = ',')
```

Linear Regression Model

- Import : `from sklearn import linear_model`

Linear Regression Model

- Import : `from sklearn import linear_model`
- Select Model
`regr = linear_model.LinearRegression(fit_intercept=True)`

Linear Regression Model

- Import : `from sklearn import linear_model`
- Select Model
`regr = linear_model.LinearRegression(fit_intercept=True)`
- Train the model using training data
`regr.fit(X_train, y_train)`
- Make predictions using test data
`y_pred = regr.predict(X_test)`

Linear Regression Model

- To retrieve coefficients: $\theta_1, \theta_2, \dots$
`regr.coef_`
- To retrieve coefficient θ_0
`regr.intercept_`

Assignment 1: Single Variable Linear Regression

Load file “data_train_sv.csv” to `train_data`

Load file “data_test_sv.csv” to `train_data`

- Train linear regression model using `train_data` and predict the `target` values of `test_data`
- Compute Mean squared error
- Plot the model

Linear Regression : Evaluation Metrics

- Import `from sklearn.metrics import mean_squared_error`

Linear Regression : Evaluation Metrics

- Import `from sklearn.metrics import mean_squared_error`
- `mean_squared_error(y_test, y_pred)`

Linear Regression

Training the model

`regr.fit(X_train, y_train)`

$$\theta = (X_{train}^T * X_{train})^{-1} * X_{train}^T * y_{train}$$

Linear Regression

Training the model

```
regr.fit(X_train, y_train)
```

$$\theta = (X_{train}^T * X_{train})^{-1} * X_{train}^T * y_{train}$$

Prediction of target value of test data

```
y_pred = regr.predict(X_test)
```

$$y_{pred} = X_{test}\theta$$

- `import matplotlib.pyplot as plt`
`plt.scatter(X_test, y_test, color='black')`

- `import matplotlib.pyplot as plt`
 `plt.scatter(X_test, y_test, color='black')`
 `plt.plot(X_test, y_pred, color='blue', linewidth=3)`

- `import matplotlib.pyplot as plt`
 `plt.scatter(X_test, y_test, color='black')`
 `plt.plot(X_test, y_pred, color='blue', linewidth=3)`
 `plt.show()`

Assignment 1: Multiple Variable Linear Regression

Load file “housing_data.csv” to **data**

Split 80% of data to training data and 20% of data to test data

Normalize features (Feature scaling)

- Train linear regression model using **train_data** and predict the **target** values of **test_data**
- Compute Mean squared error

Training and Testing Data Split

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.20, random_state=42)
```

Feature Scaling

```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
scaler.fit(X_train)  
X_train = scaler.transform(X_train)  
X_test = scaler.transform(X_test)
```


Assignment 3: Polynomial Curve Fitting

Dataset Generation

Let the underlying function be

$$y_{actual} = \cos^2 2\pi X$$

Plot the function y_{actual}

```
X = np.linspace(0, 0.5, 100)
y_actual = np.cos(2 * np.pi * X)**2
```

Assignment 3: Polynomial Curve Fitting

Dataset Generation

Let the underlying function be

$$y_{actual} = \cos^2 2\pi X$$

Plot the function y_{actual}

```
X = np.linspace(0, 0.5, 100)  
y_actual = np.cos(2 * np.pi * X)**2
```

Generate the dataset by adding noise to the underlying function

Assignment 3: Polynomial Curve Fitting

Dataset Generation

Let the underlying function be

$$y_{actual} = \cos^2 2\pi X$$

Plot the function y_{actual}

```
X = np.linspace(0, 0.5, 100)
y_actual = np.cos(2 * np.pi * X)**2
```


Generate the dataset by adding noise to the underlying function

```
noise = np.random.normal(0, 0.1, 100)
y = y_actual + noise
```

Plot underlying function y_{actual} and y


Assignment 3: Polynomial Curve Fitting

Dataset Generation



Assignment 3: Polynomial Curve Fitting


Dataset Generation



Assignment 3: Polynomial Curve Fitting

Fit linear regression model

Assignment 3: Polynomial Curve Fitting



Assignment 3: Polynomial Curve Fitting


- Import

```
from sklearn.preprocessing import PolynomialFeatures  
from sklearn.pipeline import Pipeline
```

- ① `deg=2`
- ② `polynomial_features =
PolynomialFeatures(degree=deg,include_bias=True)`
- ③ `linear_regression = linear_model.LinearRegression()`
- ④ `pipeline = Pipeline([("polynomial_features", polynomial_features),
("linear_regression", linear_regression)])`
- ⑤ `pipeline.fit(X_train, y_train)`


Assignment 3: Polynomial Curve Fitting

deg=1




Assignment 3: Polynomial Curve Fitting

deg=10




Assignment 3: Polynomial Curve Fitting


Underfitting



Overfitting



Assignment 3: Polynomial Curve Fitting



Regularization

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y(i))^2 + \lambda \theta^T \theta$$

λ is the regularization parameter

Regularization

```
model = linear_model.Ridge(alpha=0.001, fit_intercept=True)
```

To use sample datasets in Sklearn

```
from sklearn import datasets  
data = datasets.load_breast_cancer()  
X=data.data  
y=data.target
```

Logistic Regression

```
classifier = linear_model.LogisticRegression()  
classifier.fit(X_train, y_train)  
y_pred = classifier.predict(X_test)
```